

# **HATE SPEECH DETECTION USING MACHINE LEARNING**



## **A DESIGN PROJECT REPORT**

*Submitted by*

**KABILESHWARAN K**

**MELWIN A.B**

**MOHAMED FAIZUL S**

**RAVIDHARSHEN M.K**

*in partial fulfilment for the award of the*

*degree of*

**BACHELOR OF TECHNOLOGY**

*in*

**ARTIFICIAL INTELLIGENCE AND DATA SCIENCE**

**K. RAMAKRISHNAN COLLEGE OF TECHNOLOGY**

(An Autonomous Institution, affiliated to Anna University Chennai and Approved by AICTE, New Delhi)

**SAMAYAPURAM-621112**

**JUNE-2024**

**K. RAMAKRISHNAN COLLEGE OF TECHNOLOGY  
(AUTONOMOUS)  
SAMAYAPURAM-621112  
BONAFIDE CERTIFICATE**

Certified that this design project report titled “**HATE SPEECH DETECTION USING MACHINE LEARNING**” is the bonafide work of **KABILESHWARAN K (REG NO: 811721243022) MELWIN A.B (REG NO: 811721243027) MOHAMED FAIZUL S (REG NO: 811721243030) RAVIDHARSHEN M.K (811721243043)** who carried out the project under my supervision.

**SIGNATURE**

Dr.T. Avudaiappan M.E., Ph.D.,

**HEAD OF THE DEPARTMENT**

Associate Professor

Department of AI

K. Ramakrishnan College of Technology  
(Autonomous)

Samayapuram – 621112

**SIGNATURE**

Mrs. P.Jasmine Jose M.E.,

**SUPERVISOR**

Assistant Professor

Department of AI

K. Ramakrishnan College of Technology  
(Autonomous)

Samayapuram – 621112

Submitted for the viva-voce examination held on .....

**INTERNAL EXAMINER**

**EXTERNAL EXAMINER**

## **DECLARATION**

We jointly declare that the project report on “**HATE SPEECH DETECTION USING MACHINE LEARNING**” is the result of original work done by us and best of our knowledge, similar work has not been submitted to “**ANNA UNIVERSITY CHENNAI**” for the requirement of Degree of **BACHELOR OF TECHNOLOGY**. This design project report is submitted on the partial fulfilment of the requirement of the award of Degree of **BACHELOR OF TECHNOLOGY**.

**SIGNATURE**

---

**KABILESHWARAN K**

---

**MELWIN A.B**

---

**MOHAMED FAIZUL S**

---

**RAVIDHARSHEN M.K**

**PLACE : SAMAYAPURAM**

**DATE :**

## ACKNOWLEDGEMENT

It is with great pride that we express our gratitude and in - debt to our institution “**K.RAMAKRISHNAN COLLEGE OF TECHNOLOGY (AUTONOMOUS)**”, for providing us with the opportunity to do this project.

We are glad to credit honourable chairman **Dr. K. RAMAKRISHNAN, B.E.**, for having provided for the facilities during the course of our study in college.

We would like to express our sincere thanks to our beloved Executive Director **Dr.S. KUPPUSAMY, MBA., Ph.D.**, for forwarding to our project and offering adequate duration in completing our project.

We would like to thank our Principal **Dr.N. VASUDEVAN, M.E., Ph.D.**, who gave opportunity to frame the project the full satisfaction.

We whole heartily thanks to **Dr.T. AVUDAIAPPAN, M.E., Ph.D.**, HEAD OF THE DEPARTMENT, **ARTIFICIAL INTELLIGENCE** for providing his encourage pursuing this project.

I express my deep and sincere gratitude to my project guide **Mrs. P. JASMINE JOSE M.E.**, ASSISTANT PROFESSOR, **ARTIFICIAL INTELLIGENCE** for his incalculable suggestions, creativity, assistance and patience which motivated me to carry out the project successfully.

I render my sincere thanks to my project coordinator **Mrs. G.NALINA KEERTHANA ,ME.**, ,other faculties and non-teaching staff members for providing valuable information during the course. I wish to express my special thanks to the officials & Lab Technicians of our departments who rendered their help during the period of the work progress.

## **ABSTRACT**

Hate speech detection has emerged as a vital research domain within natural language processing and machine learning due to the increasing spread of harmful content on online platforms. This study delves into the creation of machine learning models aimed at the automatic identification and classification of hate speech in text. A diverse dataset sourced from various social media platforms is utilized, employing a range of machine learning techniques, including traditional methods such as Naive Bayes and Support Vector Machines, along with contemporary approaches like deep learning using recurrent neural networks and transformers. Effective feature representation of textual data is achieved through techniques such as TF-IDF and word embeddings. The study addresses challenges such as the nuanced and context-dependent nature of hate speech, including the detection of implicit and coded language. The models performance is evaluated using metrics like accuracy, precision, recall, and F1-score. Findings indicate that advanced models, particularly those based on transformer architectures, offer substantial improvements over traditional methods in the detection and classification of hate speech. This research not only enhances the academic understanding of hate speech detection but also has practical implications for the development of automated systems designed to mitigate online hate speech, there by promoting safer digital environments.

## TABLE OF CONTENTS

CHAPTER NO	TITLE	PAGENO
	ABSTRACT	V
	LIST OF FIGURES	IX
	LIST OF ABBREVIATIONS	X
1	INTRODUCTION	11
	1.1 BACKGROUND	11
	1.2 PROBLEM STATEMENT	12
	1.3 AIM & OBJECTIVE	12
	1.3.1 Aim	12
	1.3.2 Objectives	13
2	LITERATURE SURVEY	14
	2.1 DETECTING HATE SPEECH IN MULTI-MODAL MEMES.	14
	2.2 SOCIAL SHOUT-HATE SPEECH DETECTION USING MACHINE LEARNING ALGORITHM	15
	2.3 DEEP LEARNING MODELS FOR MULTILINGUAL HATE SPEECH DETECTION	16
	2.4 A LITERATURE REVIEW OF TEXTUAL HATE SPEECH DETECTION METHODS AND DATASETS.	17

	2.5 A COMPARATIVE ANALYSIS OF MACHINE LEARNING ALGORITHM	18
<b>3</b>	<b>SYSTEM ANALYSIS</b>	<b>19</b>
	3.1 EXISTING SYSTEM	19
	3.1.1 Algorithm used	19
	3.2 PROPOSED SYSTEM	22
	3.2.1 Techniques used	23
<b>4</b>	<b>SYSTEM SPECIFICATION</b>	<b>25</b>
	4.1 HARDWARE SYSTEM CONFIGURATION	25
	4.2 SOFTWARE SYSTEM CONFIGURATION	25
	4.3 SOFTWARE DESCRIPTION	25
	4.3.1 Library	25
	4.3.2 Developing Environment	28
<b>5</b>	<b>ARCHITECTURE DESIGN</b>	<b>30</b>
	5.1 SYSTEM DESIGN	30
	5.2 DATA FLOW DIAGRAM	31
	5.3 USE CASE DIAGRAM	31
	5.4 ACTIVITY DIAGRAM	32
	5.5 SEQUENCE DIAGRAM	33

<b>6</b>	<b>MODULE DESCRIPTION</b>	<b>34</b>
	6.1 MODULES	34
	6.1.1 Methodology	34
	6.1.2 Data Gathering	34
	6.1.3 Kaggle	35
	6.1.4 Data Preprocessing	36
	6.1.5 Count Vectorization	36
	6.1.6 Evaluate Model	37
<b>7</b>	<b>RESULTS AND DISCUSSION</b>	<b>39</b>
<b>8</b>	<b>CONCLUSION AND FUTURE ENHANCEMENT</b>	<b>41</b>
	8.1 CONCLUSION	41
	8.2 FUTURE ENHANCEMENTS	41
	<b>APPENDIX 1.Sample code</b>	<b>42</b>
	<b>APPENDIX 2.Screen shot</b>	<b>46</b>
	<b>REFERENCES</b>	<b>47</b>



## LIST OF FIGURES

FIGURE.NO	TITLE	PAGE.NO
5.1	SYSTEM ARCHITECTURE	30
5.2	DATA FLOW DIAGRAM	31
5.3	USE CASE DIAGRAM	31
5.4	ACTIVITY DIAGRAM	32
5.5	SEQUENCE DIAGRAM	33
7.1	HOME PAGE WEB INTERFACE	39
7.2	INTERPRETATION DIAGRAM	40

## **LIST OF ABBREVIATIONS**

<b>ML</b>	MACHINE LEARNING
<b>AI</b>	ARTIFICIAL INTELLIGENCE
<b>NLP</b>	NATURAL LANGUAGE PROCESSING
<b>LSTM</b>	LONG SHORT-TERM MEMORY
<b>POS</b>	PART OF SPEECH
<b>ROC</b>	RECEIVER OPERATING CHARACTERISTIC
<b>BOW</b>	BAG OF WORD
<b>AUC</b>	AREA UNDER THE CURVE
<b>DNN</b>	DEEP NEURAL NETWORK

# **CHAPTER 1**

## **INTRODUCTION**

Hate speech, defined as any communication that disparages a person or a group based on attributes such as race, religion, ethnicity, gender, sexual orientation, disability, or nationality, has become a pervasive issue in online discourse. The proliferation of social media platforms and digital communication channels has facilitated the rapid dissemination of hate speech, contributing to the amplification of social divisions and the erosion of civil discourse. In recent years, there has been a growing recognition of the need to address hate speech in digital spaces and to develop effective strategies for its detection and mitigation. Traditional methods of manual moderation and content filtering are often inadequate to cope with the sheer volume of user-generated content across various online platforms. This has led to an increasing reliance on automated techniques, particularly machine learning, for the identification and classification of hate speech.

### **1.1 BACKGROUND**

The proliferation of hate speech on digital platforms has become a pressing societal concern, necessitating innovative solutions for its detection and mitigation. Traditional methods of content moderation are often insufficient to cope with the scale and complexity of online hate speech. Machine learning, a subset of artificial intelligence, has emerged as a promising avenue for addressing this challenge by automating the process of identifying and classifying hateful content. By leveraging annotated datasets and computational algorithms, machine learning models can learn to discern linguistic patterns and contextual cues indicative of hate speech. However, the development of effective hate speech detection models requires addressing various technical, ethical, and societal challenges, including dataset biases, algorithmic fairness, and free speech considerations.

## **1.2 PROBLEM STATEMENT**

In today's digital landscape, the pervasive spread of hate speech poses a significant threat to online discourse, social harmony, and individual well-being. The sheer volume and diverse forms of hateful content make manual moderation impractical and insufficient for effectively addressing the issue. Existing automated methods often struggle to accurately identify hate speech due to its nuanced and context-dependent nature. Consequently, there is a pressing need for robust and scalable solutions that leverage advanced technologies like machine learning to detect and mitigate hate speech effectively. However, developing such solutions requires overcoming several challenges, including defining hate speech in a computationally tractable manner, addressing biases in training data, ensuring model fairness and transparency, and navigating the complex interplay between freedom of expression and the prevention of harm. Addressing these challenges is crucial for advancing the field of hate speech detection and fostering safer online environments conducive to healthy discourse and mutual respect.

## **1.3 AIMS AND OBJECTIVES**

### **1.3.1 AIM**

- Develop algorithms capable of accurately identifying and categorizing instances of hate speech within textual data.
- Improve the safety and well-being of online communities by detecting and mitigating the spread of hateful content on digital platforms.
- Reduce the harm caused by hate speech, including discrimination, harassment, and incitement to violence, by proactively identifying and addressing such content.
- Empower users to report and address instances of hate speech by providing tools and resources for identifying and reporting offensive content.

### **1.3.2 OBJECTIVES**

- Create machine learning models capable of accurately distinguishing between hate speech and non-hate speech content with high precision and recall.
- Enhance the ability of hate speech detection models to generalize across different linguistic contexts, demographics, and online platforms.
- Extend hate speech detection capabilities to include multimodal content such as images, videos, and audio recordings, in addition to textual data.
- Identify and classify subtle forms of hate speech, including implicit biases, dog whistles, and coded language, which may evade traditional detection methods.
- Develop models capable of detecting hate speech in multiple languages to address the global nature of online communication and hate speech dissemination.
- Identify and mitigate biases in hate speech detection models to ensure fair and equitable treatment across diverse demographic groups and cultural contexts.
- Implement algorithms capable of detecting hate speech in real-time to facilitate timely intervention and response by platform moderators and administrators.
- Design scalable hate speech detection systems capable of processing large volumes of user-generated content on popular online platforms with minimal computational resources.
- Enhance the interpretability and transparency of hate speech detection models by providing insights into the features and decision-making processes underlying classification results.
- Incorporate ethical considerations into hate speech detection algorithms, including privacy protection, freedom of expression, and the prevention of unintended consequences or harm.

## CHAPTER 2

### LITERATURE SURVEY

**2.1 TITLE:** Detecting Hate Speech in multi-modal Memes

**AUTHOR:** Abhishek Das, Japsimar Singh Wahi, Siyao Li

**YEAR OF PUBLICATION:** 2020

**ALGORITHM USED:** For the proposed approach of detecting hate speech in multi-modal memes, a combination of natural language processing (NLP) algorithms and computer vision techniques is employed. In the textual domain, algorithms like Word Embeddings and Text Classification Models are utilized to extract features from meme captions or associated text and classify them as hate speech or non-hate speech. Simultaneously, in the visual domain, Convolutional Neural Networks (CNNs) and pre-trained models such as VGG or ResNet are employed to extract visual features from meme images.

**ABSTRACT:** Detecting hate speech in multi-modal memes presents a complex challenge due to the fusion of textual and visual elements, often laden with implicit meanings and cultural context. This paper proposes a novel approach leveraging machine learning techniques to analyze both textual and visual features for hate speech detection in memes.

**MERITS:** Integrates both textual and visual features for hate speech detection, capturing the nuanced nature of multi-modal memes.

**DEMERITS:** Integrating both text and image analysis techniques may increase computational overhead, requiring significant resources for processing large volumes of multi-modal memes.

**2.2 TITLE:** Social Shout -Hate Speech Detection Using Machine Learning Algorithm.

**AUTHOR:** V.B. Ohol, Siddhi Patil, Ishwari Gamne, Sayali Patil, Shweta Bandawane.

**YEAR OF PUBLICATION:** 2023.

**ALGORITHM USED:** The project preprocesses text data, selects a machine learning algorithm (ranging from traditional to deep learning models), trains it on labeled data to detect hate speech patterns, and evaluates its performance using accuracy metrics, aiding in combating online toxicity.

**ABSTRACT:** Leveraging techniques from natural language processing and machine learning, this project explores various algorithms and features to effectively distinguish between hate speech and non-hate speech messages. Through extensive experimentation and evaluation on labeled datasets, the proposed approach demonstrates promising results in accurately detecting hate speech, thereby contributing to efforts in combating online toxicity and promoting safer online communities.

**MERITS:** Firstly, by automating the identification process, it reduces reliance on manual moderation, enhancing efficiency and scalability. Secondly, its flexibility allows for customization to specific platforms or communities, ensuring relevance and accuracy. Additionally, continuous learning through model updates enables adaptation to evolving language patterns, improving effectiveness over time.

**DEMERITS:** Biases from training data, misclassifications, continuous adaptation needs, and computational demands pose challenges to hate speech detection algorithms.

**2.3 TITLE:** Deep Learning Models For Multilingual Hate Speech Detection.

**AUTHOR:** Sai Saketh Aluru, Binny Mathew, Punyajoy Saha, Animesh Mukherjee.

**YEAR OF PUBLICATION:** 2022.

**ALGORITHM USED:** Deep Learning models for multilingual hate speech detection often utilize architectures such as Convolutional Neural Networks (CNNs), Recurrent Neural Networks (RNNs), Long Short-Term Memory (LSTM) networks, Gated Recurrent Units (GRUs), and Transformer-based architectures like BERT (Bidirectional Encoder Representations from Transformers).

**ABSTRACT:** Multilingual hate speech detection using deep learning involves training models on datasets containing text samples in multiple languages annotated for hate speech. These models learn representations of language that capture semantic and syntactic cues indicative of hate speech. The trained models can then classify new text samples into hate speech or non-hate speech categories, thus aiding in moderating online content across different languages.

**MERITS:** Deep learning models can learn representations of language that generalize across different languages, enabling effective hate speech detection in multilingual contexts. Deep learning models can learn complex features directly from raw text data, eliminating the need for manual feature engineering.

**DEMERITS:** Deep learning models for hate speech detection require large amounts of annotated data in multiple languages, which may be challenging and expensive to obtain.



**2.4 TITLE:** A Literature Review Of Textual Hate Speech Detection Methods And Datasets

**AUTHOR:** Fatimah Alkomah, Xiaogang Ma

**YEAR OF PUBLICATION:** 2022

**ALGORITHM USED:** The literature review likely covers a range of algorithms used for textual hate speech detection, including traditional machine learning algorithms like Support Vector Machines (SVM), Logistic Regression, Naive Bayes, as well as deep learning architectures such as Convolutional Neural Networks (CNNs), Recurrent Neural Networks (RNNs), Long Short-Term Memory (LSTM) networks, Gated Recurrent Units (GRUs), and Transformer-based models like BERT (Bidirectional Encoder Representations from Transformers).

**ABSTRACT:** It synthesizes existing research to identify trends, challenges, and future directions in this domain. The review likely includes an analysis of various algorithms, feature representations, dataset characteristics, evaluation metrics, and advancements in hate speech detection techniques.

**MERITS:** The literature review offers a comprehensive summary of the state-of-the-art methods and datasets for textual hate speech detection, providing valuable insights for researchers and practitioners.

**DEMERIT:** The literature review may be subject to publication bias, as it might predominantly include studies that have been published in reputable journals or conferences, potentially overlooking relevant work in gray literature or unpublished sources.

**2.5 TITLE:** A comparative analysis of machine learning algorithms for hate speech detection in social media

**AUTHOR:** Esraa Omran , Estabraq Al Tararwah ,Jamal Al Qundus

**YEAR OF PUBLICATION:** 2023.

**ALGORITHM USED:** In our comparative analysis of machine learning algorithms for hate speech detection in social media, we employed a diverse set of methodologies to evaluate their effectiveness. We began with the Naive Bayes Classifier, leveraging its simplicity, speed, and efficiency, particularly beneficial for smaller datasets. Support Vector Machines (SVM) were also incorporated, known for their efficacy in high-dimensional spaces and memory efficiency, although sensitive to kernel choice and computationally demanding for larger datasets.

**ABSTRACT:** Hate speech detection in social media is a crucial task to maintain a healthy online community and mitigate the spread of harmful content. This study aims to compare various machine learning algorithms for hate speech detection. We evaluate the performance of these algorithms based on their accuracy, efficiency, and generalizability across different social media platforms.

**MERITS:** Simple, fast, and efficient with good performance on small datasets, Effective in high-dimensional spaces, memory-efficient, and versatile.

**DEMERITS:** Assumes independence among features, may not perform well with correlated features.

## **CHAPTER 3**

### **SYSTEM ANALYSIS**

#### **3.1 EXISTING SYSTEM**

Several existing systems utilize machine learning for hate speech detection, leveraging various algorithms and techniques to address this pervasive issue. These systems typically employ natural language processing (NLP) models trained on annotated datasets to classify text-based content as either hateful or benign. Some systems focus on specific platforms such as social media networks, while others offer broader coverage across online forums and websites. Key components of these systems include feature extraction methods, sentiment analysis, and context-aware classification to accurately identify hate speech while minimizing false positives.

##### **3.1.1 ALGORITHM USED**

###### **1. Convolutional neural network**

In deep learning, a convolutional neural network is a class of artificial neural networks most applied to analyze visual imagery. CNNs use a mathematical operation called convolution in place of general matrix multiplication in at least one of their layers.

###### **2.LSTM**

- LSTM stands for long short-term memory networks, used in the field of Deep Learning. It is a variety of recurrent neural networks (RNNs) that are capable of learning long-term dependencies, especially in sequence prediction problems.
- LSTM stands for Long Short-Term Memory, and it is a type of artificial recurrent neural network (RNN) architecture that is commonly used in natural language processing (NLP) tasks such as text classification, language translation and speech recognition.

LSTM networks are designed to address the vanishing gradient problem that is often encountered in standard RNNs. The vanishing gradient problem occurs when gradients

propagated back through the network during training become too small to update the weights effectively. This problem can make it difficult for RNNs to learn long-term dependencies in sequential data. LSTM networks use a more complex architecture than standard RNNs, with additional memory cells and gates that control the flow of information through the network. The three key components of an LSTM network are:

1. Cell state: This is the memory of the LSTM network. The cell state is updated by adding or removing information from it via gates.
2. Gates: Gates are used to control the flow of information through the network. They are made up of a sigmoid neural net layer and a pointwise multiplication operation. The three types of gates in an LSTM network are the input gate, the forget gate, and the output gate.
3. Hidden state: This is the output of the LSTM network, which is based on the cell state and the input at the current time step.

### **Bi-directional lstm:**

A Bidirectional LSTM, or bi LSTM, is a sequence processing model that consists of two LSTMs: one taking the input in a forward direction, and the other in a backwards direction.

### **XGBoost:**

XGBoost (extreme Gradient Boosting) is a supervised machine learning algorithm that is used for both classification and regression tasks. It is an ensemble method that combines multiple weak learners, typically Decision trees, to create a more accurate model. XGBoost is a fast and efficient algorithm and has been used by the winners of many data science competitions. XGBoost works only with numeric variables and have already replaced the categorical variables with numeric variables.

### **3. Stochastic gradient descent algorithm**

Stochastic Gradient Descent (SGD) is a variant of the Gradient Descent algorithm used for optimizing machine learning models. In this variant, only one random training example is used to calculate the gradient and update the parameters at each iteration. Here are some of the advantages and disadvantages of using SGD. Suppose, you have a million samples in

your dataset, so if you use a typical Gradient Descent optimization technique, you will have to use all of the one million samples for completing one iteration while performing the Gradient Descent, and it has to be done for every iteration until the minima are reached. Hence, it becomes computationally very expensive to perform. This problem is solved by Stochastic Gradient Descent. In SGD, it uses only a single sample, i.e., a batch size of one, to perform each iteration. The sample is randomly shuffled and selected for performing the iteration. Stochastic gradient descent (SGD) is an iterative optimization algorithm commonly used in machine learning to minimize the cost function of a model. It is a variant of gradient descent that updates the model's parameters based on a randomly selected subset of training examples, known as a mini-batch, instead of the entire dataset. The algorithm works by first initializing the model's parameters with random values. Then, for each iteration, the algorithm randomly selects a mini-batch of training examples from the dataset and computes the gradient of the cost function with respect to the parameters using only the examples in the mini-batch. The algorithm then updates the parameters in the opposite direction of the gradient, scaled by a learning rate parameter. The learning rate determines the step size of the update and must be chosen carefully to balance between convergence speed and stability. The process is repeated for a fixed number of iterations or until convergence, which is typically determined by monitoring the change in the cost function or the model's performance on a validation set. The SGD algorithm is widely used in deep learning, where the datasets are often large and computationally expensive to process, making it impractical to compute the gradient over the entire dataset in each iteration. Let Discuss detail in further section.

### **3.2 PROPOSED SYSTEM**

Our proposed hate speech detection system introduces a novel approach that combines the strengths of machine learning with emerging techniques in sociolinguistics and cognitive psychology. Unlike traditional systems that primarily rely on textual analysis, our system incorporates multimodal data sources, including text, images, and audio, to provide a more comprehensive understanding of hate speech phenomena. By integrating multimodal analysis, our system can capture nuanced expressions of hate speech that may manifest through different modalities, enhancing detection accuracy and contextual comprehension. Furthermore, our system emphasizes the importance of contextual relevance and cultural sensitivity in hate speech detection. We integrate domain-specific knowledge bases and ontologies to contextualize language use within specific cultural and social contexts, reducing the risk of misinterpretation or bias in classification. Additionally, our system employs advanced sociolinguistic models to analyze socio-cultural factors such as power dynamics, identity politics, and group dynamics, providing deeper insights into the underlying motivations and implications of hate speech. To address the dynamic nature of hate speech and linguistic evolution, our system adopts a continuous learning framework. By leveraging active learning strategies and reinforcement learning algorithms, our system iteratively improves its performance over time, adapting to evolving linguistic trends and emerging forms of hate speech. Moreover, we prioritize user engagement and collaboration by incorporating crowdsourcing mechanisms for data annotation and validation, fostering a community-driven approach to hate speech detection. In summary, our proposed hate speech detection system represents a significant advancement in the field by integrating multimodal analysis, sociolinguistic insights, and continuous learning mechanisms. By embracing interdisciplinary perspectives and leveraging cutting-edge technologies, our system aims to enhance the effectiveness and accuracy of hate speech detection,

ultimately contributing to the promotion of tolerance, diversity, and inclusivity in online discourse.

### 3.2.1 TECHNIQUES USED

Several algorithms are commonly used in hate speech detection using machine learning. These algorithms are applied in various stages of the detection process, including data preprocessing, feature extraction, model training, and classification. Here are some of the key algorithms used:

**1. Naive bayes classifier:** Naive Bayes classifiers are simple probabilistic classifiers based on Bayes' theorem with the assumption of independence between features. They are often used for text classification tasks, including hate speech detection, due to their simplicity and efficiency.

**2.Support vector machines (SVM):** SVM is a supervised learning algorithm that is effective for binary classification tasks. SVM aims to find the hyperplane that best separates the data points of different classes. It has been widely used in hate speech detection for its ability to handle high-dimensional data and find complex decision boundaries.

**3.Logistic regression:** Logistic regression is a statistical model used for binary classification. It models the probability of a binary outcome based on one or more predictor variables. In hate speech detection, logistic regression models are trained on textual features extracted from the input data.

**4. Decision trees and random forests:** Decision trees are hierarchical tree structures used for classification tasks. Random forests are ensembles of decision trees trained on random subsets of the data. Both decision trees and random forests are used in hate speech detection for their interpretability and ability to capture complex decision boundaries.

**5. Recurrent neural networks (RNNs):** RNNs are a class of neural networks designed to process sequential data. They are often used for natural language processing tasks, including hate speech detection, due to their ability to capture dependencies between words in text data.

**6. Convolutional neural networks (CNNs):** CNNs are neural networks designed to process grid-like data, such as images or text. In hate speech detection, CNNs can be applied to text data by treating words or characters as pixels in an image, enabling the detection of patterns and features at different levels of abstraction.

**7.Long short-term memory (LSTM) Networks:** LSTMs are a type of RNN that are capable of learning long-term dependencies in sequential data. They have been successfully applied to hate speech detection tasks, particularly when dealing with longer textual inputs or capturing semantic relationships between words.

**8.BERT (Bidirectional Encoder Representations from Transformers):** BERT is a transformer-based model that has achieved state-of-the-art performance on various natural language processing tasks, including hate speech detection. BERT-based models can effectively capture contextual information and semantic relationships in text data, leading to improved detection accuracy.

These algorithms can be used individually or in combination with each other, depending on the specific requirements of the hate speech detection task and the characteristics of the input data. Additionally, techniques such as ensemble learning, transfer learning, and active learning can be employed to further enhance the performance of hate speech detection systems.



## **CHAPTER 4**

### **SYSTEM SPECIFICATION**

#### **4.1 HARDWARE SYSTEM CONFIGURATION**

- The Software is a set of instructions that are used to command any systems to perform any operations . The software has the advantage to make decisions and to hardware sensible results and is useful in handling complex situations.
- OS : Windows 10, Linux
- LANGUAGE : Python
- LIBRARIES : Pandas, NumPy, Matplotlib, Seaborn.

#### **4.2 SOFTWARE SYSTEM CONFIGURATION**

- Python programming language - Python 3.x installed on the computer/server
- Operating system - Windows, Linux, or macOS.
- Python libraries such as – opencv , numpy, pandas, tensorflow, Scikit-Learn.
- Django or Flask (Python): For developing the backend of a web-based traffic prediction dashboard.
- HTML/CSS or JavaScript - for UI design

#### **4.3 SOFTWARE DESCRIPTION**

Hate Speech Detection Software is a powerful tool designed to identify and mitigate hate speech within digital content using advanced machine learning techniques. Hate speech, characterized by offensive language, discriminatory remarks, and expressions of hostility towards marginalized groups, poses a significant challenge in online platforms, social media, and various digital communication channels.

##### **4.3.1 LIBRARY**

## **1. Visual studio code ide :**

Visual Studio Code is a freeware source-code editor made by Microsoft for Windows, Linux and MacOS. Features include support for debugging, syntax highlighting, intelligent code completion, snippets, code refactoring, and embedded Git. Users can change the theme, keyboard shortcuts, preferences, and install extensions that add additional functionality. In the Stack Overflow 2019 Developer Survey, Visual Studio Code was ranked the most popular developer environment tool, with 50.7% of 87,317 respondents reporting that they use it. VS Code can be extended via extensions, available through a central repository. This includes additions to the editor and language support. A notable feature is the ability to create extensions that add support for new languages, themes, and debuggers, perform static code analysis, and add code linters using the Language Server Protocol. Visual Studio Code includes multiple extensions for FTP, allowing the software to be used as a free alternative for web development.

## **2. Python (programming language):**

Python is an interpreted, high-level, and general-purpose programming language. Python's design philosophy emphasizes code readability with its notable use of significant indentation. Its language constructs and object-oriented approach aim to help programmers write clear, logical code for small and large-scale projects. Python is dynamically-typed and garbage-collected. It supports multiple programming paradigms, including structured (particularly, procedural), object-oriented and functional programming. Python is often described as a "batteries included" language due to its comprehensive standard library. Python was designed to be highly extensible (with modules). This compact modularity has made it particularly popular as a means of adding programmable interfaces to existing applications. Libraries such as NumPy matplotlib allow the effective use of Python in scientific computing, with specialized libraries such as Bio python and Atrophy providing domain-specific functionality. Sage Math is a mathematical software with a notebook interface programmable

in Python: its library covers many aspects of mathematics, including algebra, combinatorics, numerical mathematics, number theory, and calculus. Open CV has python bindings with a rich set of features for computer vision and image processing. Python is commonly used in artificial intelligence projects and machine learning projects with the help of libraries like TensorFlow, Keras, Pytorch and Sickler. As a scripting language with modular architecture, simple syntax and rich text processing tools, Python is often used for natural language processing.

### **3.Pandas:**

Pandas is an open-source library that is made mainly for working with relational or labeled data both easily and intuitively. It provides various data structures and operations for manipulating numerical data and time series. Pandas is an open-source Python library that is widely used for data analysis and manipulation. It provides powerful data structures for working with structured data, such as tabular data, time series, and matrix data. Pandas is built on top of the NumPy library and is designed to be fast and efficient. The main data structures provided by Pandas are the Series and DataFrame. A Series is a one-dimensional array-like object that can hold any data type, while a DataFrame is a two-dimensional table of data with rows and columns. DataFrames are particularly useful for working with tabular data, as they allow for easy indexing, grouping, and aggregation.

Some of the key features of Pandas include:

- 1.Data reading and writing: Pandas can read and write data from a wide range of formats, including CSV, Excel, SQL databases, and JSON.
- 2.Data manipulation: Pandas provides a wide range of functions for manipulating and cleaning data, including filtering, grouping, merging, and reshaping.
- 3.Data visualization: Pandas integrates with popular visualization libraries like Matplotlib and Seaborn to provide powerful visualization tools for exploring and understanding data.
- 4.Missing data handling: Pandas provides tools for handling missing data, such as filling in missing values or dropping rows or columns with missing data.

5. Time series analysis: Pandas provides tools for working with time series data, including date and time parsing, resampling, and shifting. Overall, Pandas is a versatile and powerful library that is widely used in data analysis and machine learning workflows.

### **4.3.2 Developing environment**

Hate speech detection typically works by using natural language processing (NLP) techniques to analyze text and identify specific linguistic features that are associated with hate speech. Here is a general overview of how it works:

#### **Data collection**

The first step in hate speech detection is to collect a dataset of text that has been labeled as either hateful or non-hateful. This dataset is then used to train the machine learning model.

#### **Preprocessing**

The next step is to preprocess the text data by cleaning and normalizing it. This involves removing any irrelevant or noisy data, such as punctuation or special characters, and transforming the text into a standard format that the machine learning model can understand.

#### **Feature extraction**

In this step, NLP techniques are used to extract specific linguistic features that are associated with hate speech. These features can include things like profanity, derogatory language, offensive terms and slurs, and threats or incitement to violence.

#### **Machine learning**

Once the features have been extracted, they are used to train a machine learning model to classify text as either hateful or non-hateful. Various machine learning algorithms can be used for this task, including logistic regression, support vector machines (SVMs), and deep learning neural networks.

## **Evaluation**

After the machine learning model has been trained, it is evaluated using a test dataset to determine its accuracy and performance. The model is typically evaluated based on metrics such as precision, recall, and F1 score.

## **Deployment**

Once the machine learning model has been trained and evaluated, it can be deployed to classify new text data as either hateful or non-hateful. This can be done in real-time, such as for online moderation of social media platforms, or as part of a batch process for analyzing large volumes of text data.

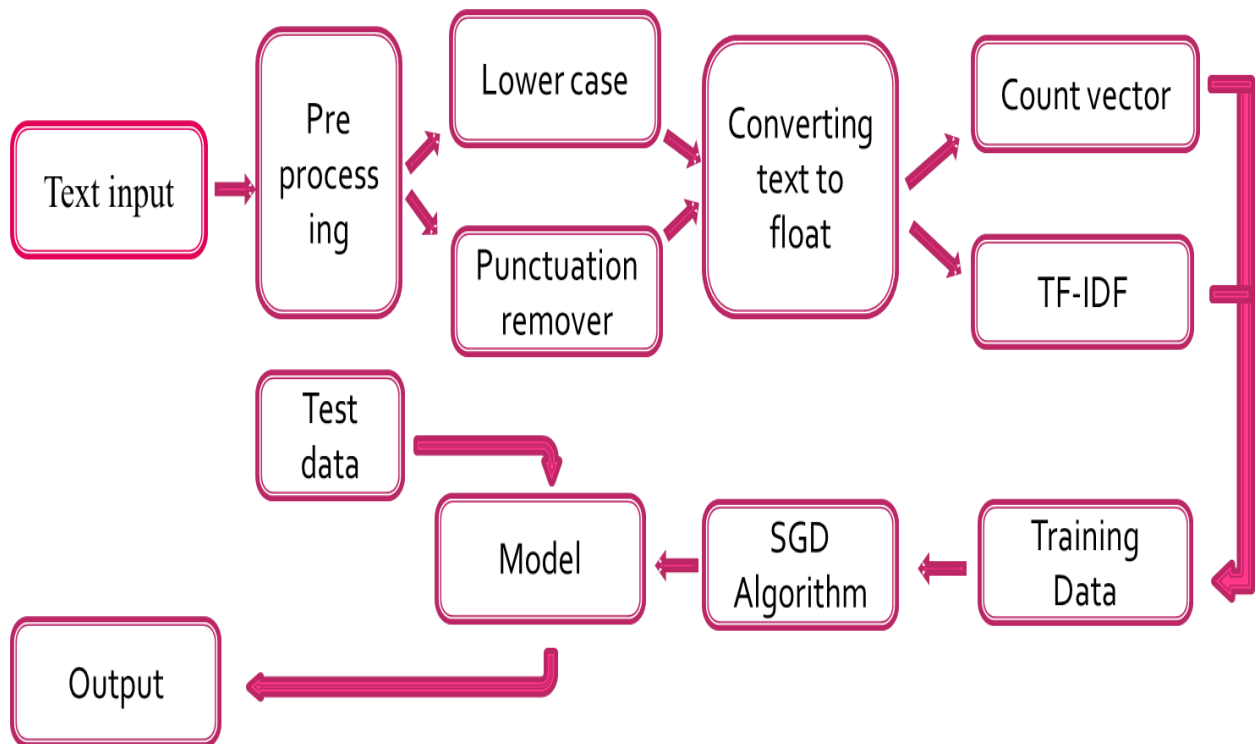
Overall, hate speech detection uses a combination of NLP techniques and machine learning algorithms to analyze text and identify specific linguistic features that are associated with hate speech. By accurately detecting hate speech, these models can help promote a safer and more inclusive online environment.

## CHAPTER 5

### ARCHITECTURAL DESIGN

#### 5.1 SYSTEM DESIGN

The hate speech detection system preprocesses text by converting it to lower case and removing punctuation. The text is then transformed into numerical representations using Count Vectorization or TF-IDF. A model, trained using the Stochastic Gradient Descent (SGD) algorithm on this data, predicts hate speech, providing output based on new text input.



**FIG:5.1 SYSTEM DESIGN**

## 5.2 DATA FLOW DIAGRAM

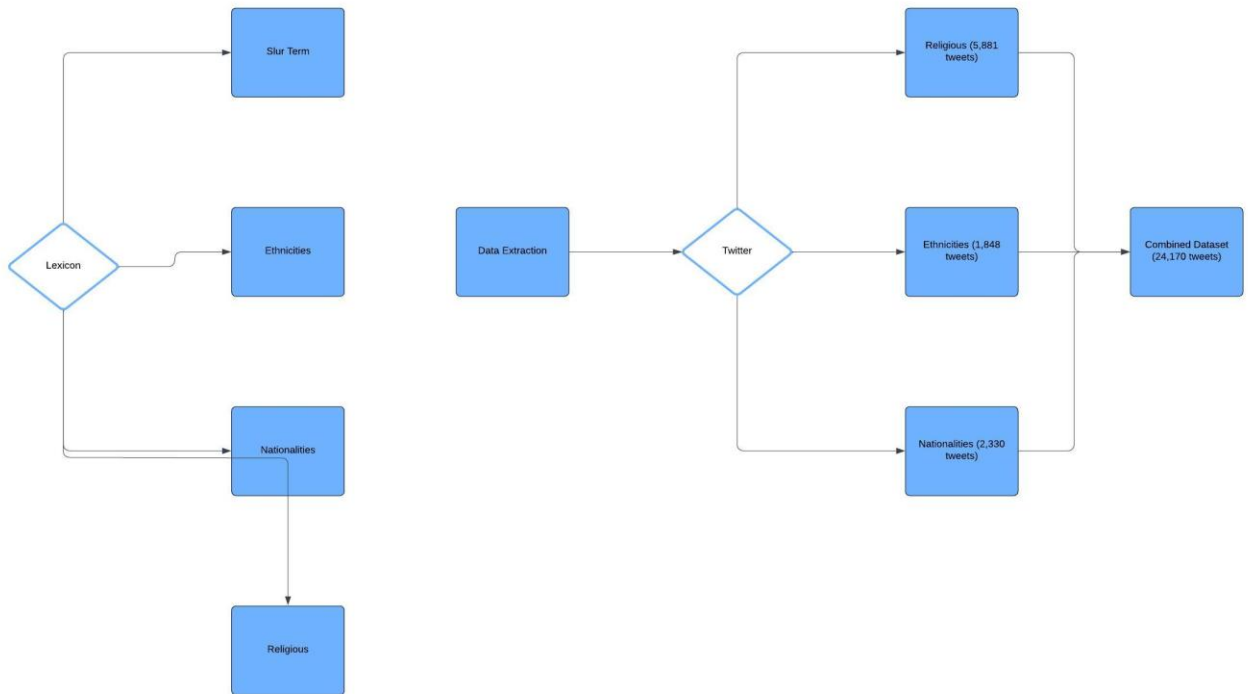


FIG:5.2 DATA FLOW DIAGRAM

## 5.3 USE CASE DIAGRAM

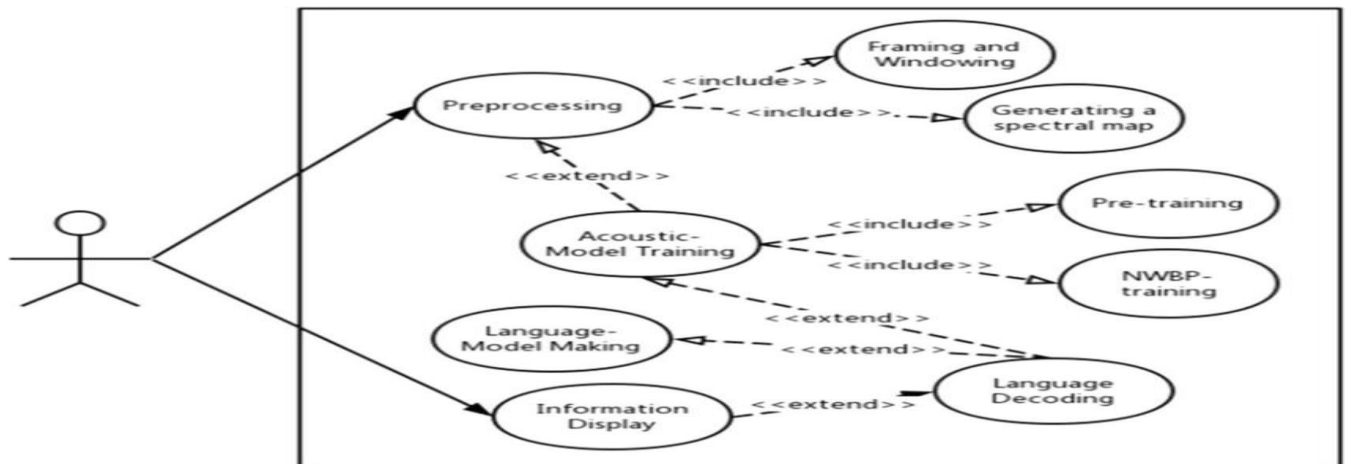
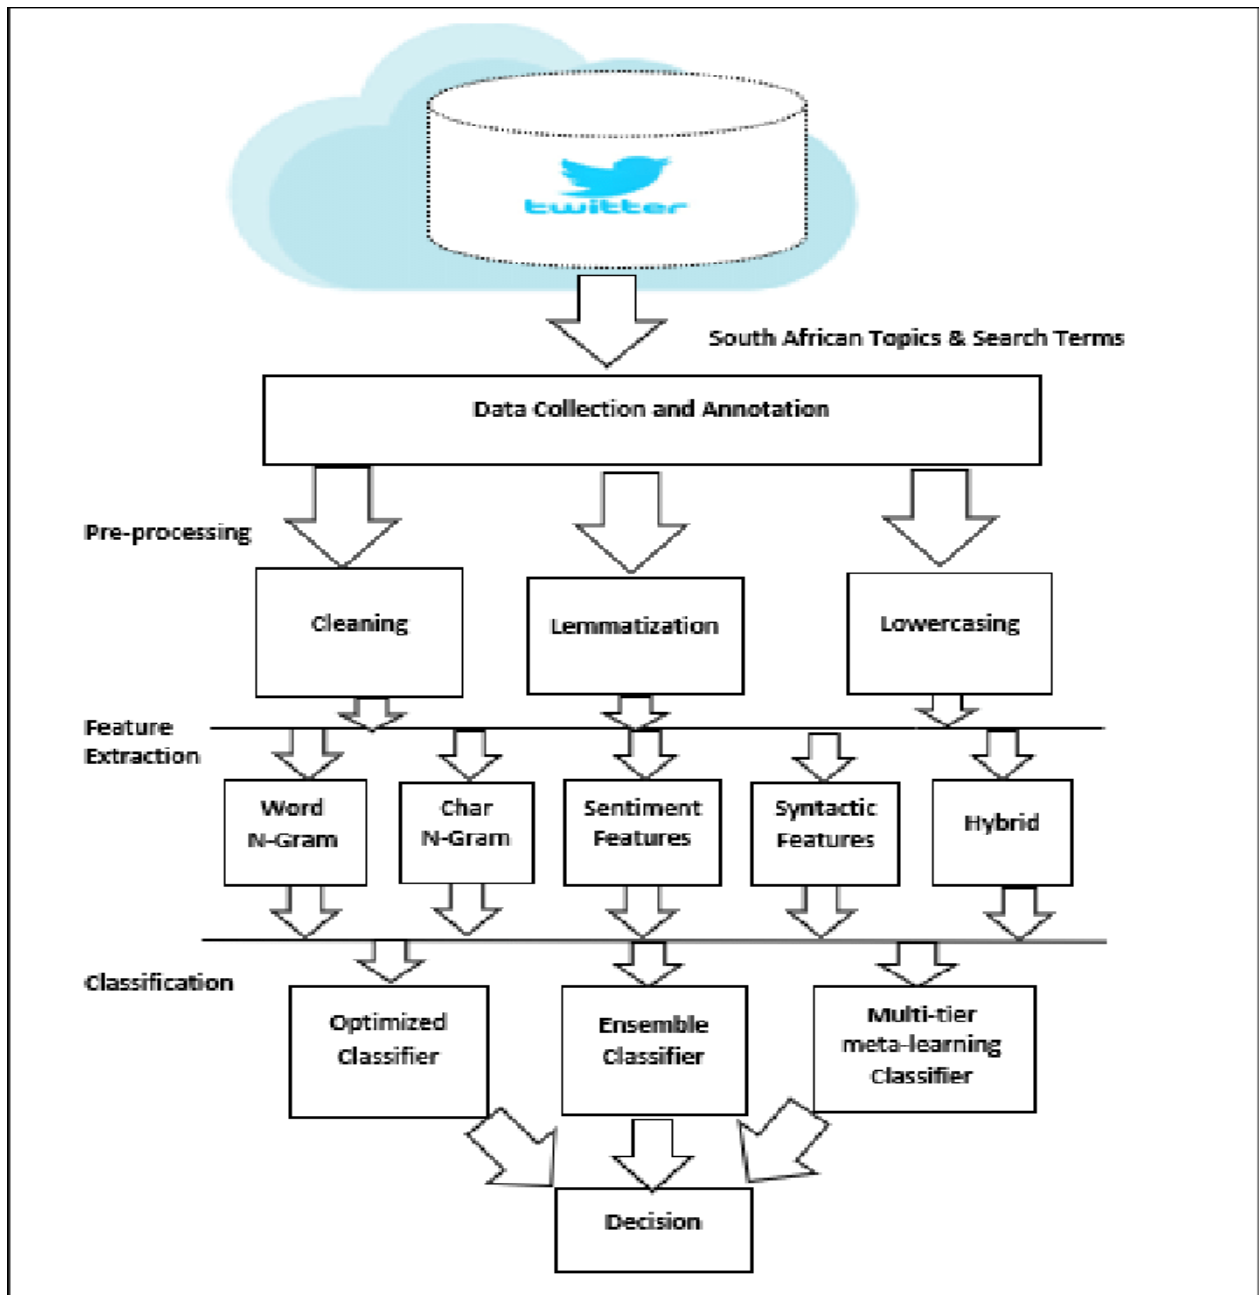


FIG:5.3 USE CASE DIAGRAM

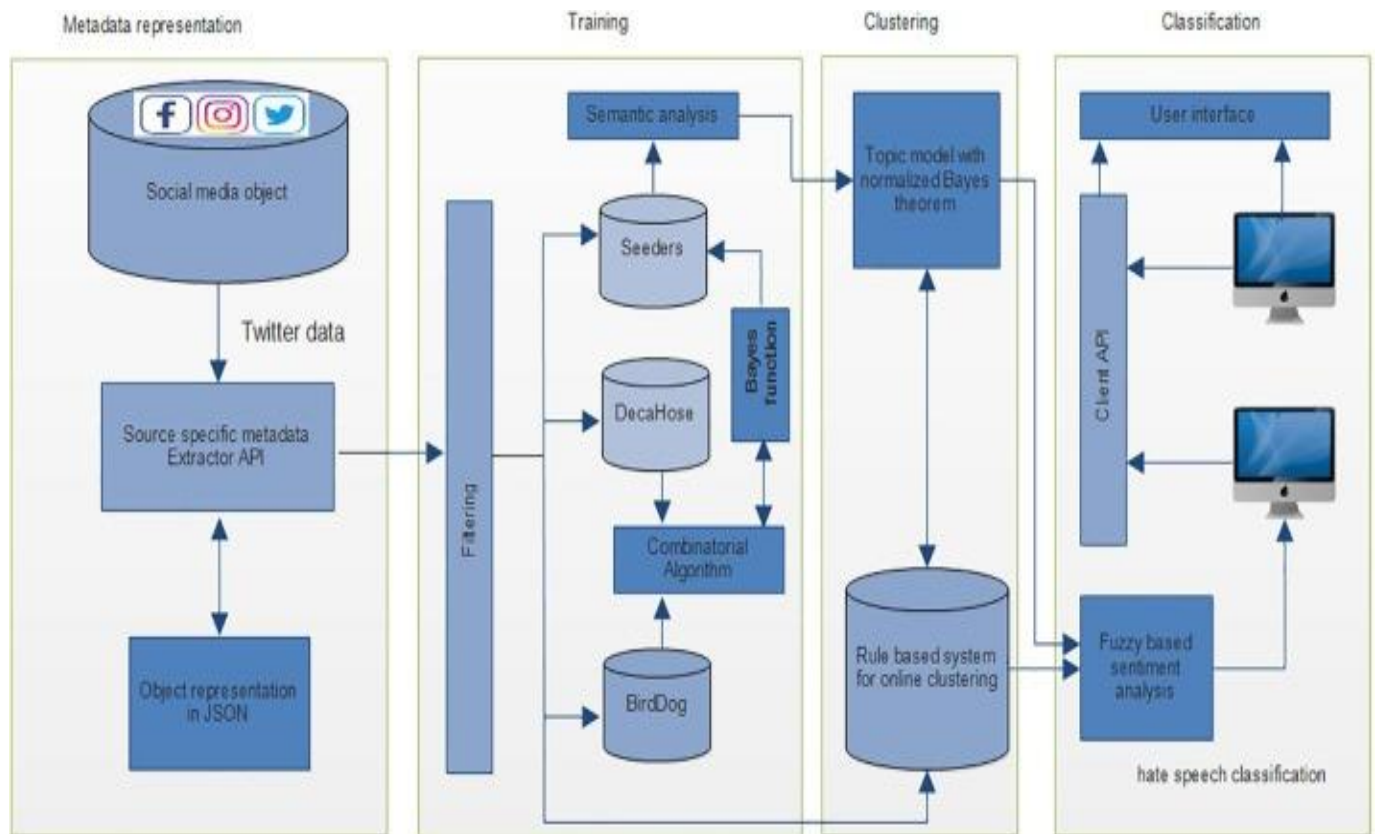
## 5.4 ACTIVITY DIAGRAM



**FIG:5.4 ACTIVITY DIAGRAM**



## 5.5 SEQUENCE DIAGRAM



**FIG:5.5 SEQUENCE DIAGRAM**

## **CHAPTER 6**

### **MODULE DESCRIPTION**

#### **6.1 MODULES**

##### **LIST OF MODULES**

- Data gathering
- Data Preprocessing
- Precision, recall
- Bot creating, prediction

##### **6.1.1 METHODOLOGY**

The methodology of our proposed project has Five major steps that involves.

- ❖Dataset collection,
- ❖Data preprocessing,
- ❖Features extraction,
- ❖Classification,
- ❖Model training and NLP.

##### **6.1.2 DATA GATHERING**

The Quant Connect Data Explorer is a portal to directly manipulate, validate and explore the Quant Connect back testing data source taken a radically open approach. Let's work together to provide the world's first crowd- curated data library for the community.

### 6.1.3 KAGGLE

Kaggle is a popular online platform for data scientists and machine learning practitioners to find and participate in data science competitions, as well as to access datasets and explore data-related resources. Here are some general steps to follow when gathering data from Kaggle.

1. Create an account on Kaggle: You will need to sign up for a Kaggle account before you can access the datasets.
2. Browse the Kaggle datasets: Kaggle offers a variety of datasets, including datasets related to machine learning, natural language processing, computer vision, and more. You can browse the datasets by topic or search for specific datasets using keywords.
3. Choose a dataset: Once you find a dataset that interests you, read the description and any accompanying documentation to ensure that it meets your needs. Check the license and terms of use to make sure you can use the data for your intended purposes.
4. Download the dataset: Kaggle provides different formats to download the datasets, you can choose between CSV, JSON, SQLite, etc.
5. Understand the data: Before you start working with the data, it's important to understand its structure and any quirks it might have. This includes identifying missing values, outliers, and other issues that could affect your analysis.
6. Prepare the data: Depending on the data and your analysis goals, you may need to preprocess the data before you can use it. This can include cleaning the data, performing feature engineering, and normalizing or scaling the data.

Overall, Kaggle can be a great resource for gathering datasets, but it's important to carefully evaluate the quality and suitability of the data before using it for any analysis or modeling.

#### **6.1.4 DATA PREPROCESSING**

Data preprocessing is a process of preparing the raw data and making it suitable for a machine learning model. It is the first and crucial step while creating a machine learning model. When creating a machine learning project, it is not always the case that we come across clean and formatted data. And while doing any operation with data, it is mandatory to clean it and put it in a formatted way. So for this, we use data preprocessing tasks. TF-IDF stands for Term Frequency-Inverse Document Frequency. It is a numerical statistic that is used to represent the importance of a term in a document or a collection of documents. The TF-IDF score of a term in a document is calculated by multiplying the term frequency (TF) and the inverse document frequency (IDF). The term frequency is simply the number of times a term appears in a document, while the inverse document frequency is a measure of how rare or common a term is across all documents in a collection. The higher the TF-IDF score of a term in a document, the more important the term is to that document. Terms with higher TF-IDF scores are more likely to be useful for information retrieval, text classification.

#### **6.1.5 COUNT VECTORIZATION**

Count Vectorization is a technique used to transform text data into a numerical representation that can be used in machine learning algorithms. It involves converting a collection of text documents into a matrix of word counts. In count vectorization, each document is represented as a vector of word counts, where each element in the vector represents the number of times a particular word occurs in the document. The resulting matrix, called a count matrix, can be used as input to machine learning algorithms.

The process of count vectorization involves several steps:

- **Tokenization:** The text data is first divided into individual words or tokens.
- **Vocabulary creation:** A vocabulary is created from the unique tokens in the text data. This vocabulary serves as the columns of the resulting count matrix.

- Count matrix creation: The count matrix is created by counting the number of times each token appears in each document.

Count vectorization is a simple yet powerful technique for processing text data for machine learning. It is often used as a baseline approach for natural language processing tasks and can be improved upon using more advanced techniques such as TF-IDF or word embeddings.

### **1.Punctuation Remover**

Punctuation remover is a simple text preprocessing technique that involves removing all punctuation marks from a text document. Punctuation marks are characters such as periods, commas, question marks, and exclamation points that are used to separate sentences or clarify the meaning of a sentence.

Punctuation removal is often performed as a preprocessing step before performing natural language processing tasks such as text classification, sentiment analysis, or text summarization. Removing punctuation can help simplify the text and reduce the size of the vocabulary, which can improve the performance of the machine learning algorithms used in these tasks.

## **6.1.6 EVALUATE MODEL**

### **PRECISION**

Precision is a performance metric used in machine learning and statistics to evaluate the accuracy of a model's positive predictions. It is defined as the ratio of true positives to the sum of true positives and false positives. In other words, precision measures how many of the positive predictions made by the model are correct.

Precision is particularly useful when the cost of a false positive prediction is high. For example, in a medical diagnosis task, a false positive result could lead to unnecessary medical procedures or treatments, which can be costly and potentially harmful to the patient.

## **RECALL**

Recall is a performance metric used in machine learning and statistics to evaluate the ability of a model to correctly identify positive instances in the data. It is defined as the ratio of true positives to the sum of true positives and false negatives. In other words, recall measures how many of the actual positive instances in the data were correctly identified by the model. Recall is particularly useful when the cost of a false negative prediction is high. For example, in a disease detection task, a false negative result could mean that a patient with the disease is not identified and does not receive the necessary treatment.

## **F-1 SCORE**

The F1 score is a harmonic mean of precision and recall, and is often used as a summary metric to evaluate the performance of a binary classification model. The F1 score combines both precision and recall into a single metric that provides a balanced evaluation of the model's performance.

Here's the formula for F1 score:

$$\text{F1 Score} = 2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$$

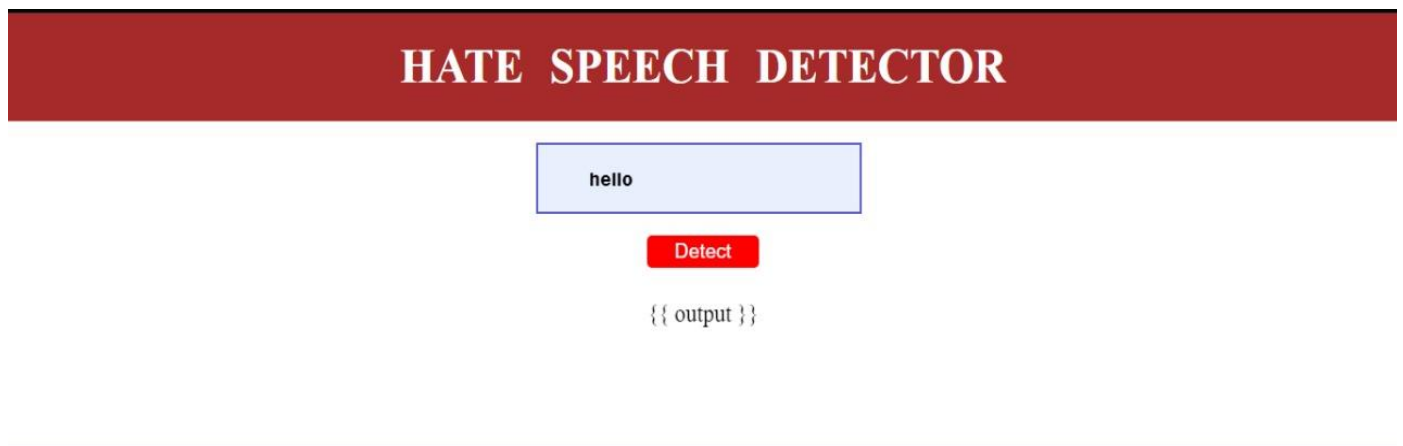
The F1 score ranges from 0 to 1, with a higher score indicating better performance. A perfect F1 score of 1 means that the model has perfect precision and recall, and is correctly identifying all the positive instances in the data.

The F1 score is a useful metric when both precision and recall are important for a given task. For example, in a fraud detection task, both precision and recall are important to minimize the number of false positives and false negatives, respectively. In this case, a high F1 score would indicate that the model is effectively balancing both precision and recall. It's worth noting that the F1 score can be biased towards precision or recall, depending on the distribution of the data and the threshold used to make predictions. In some cases, other metrics such as the area under the ROC curve.

## CHAPTER-7

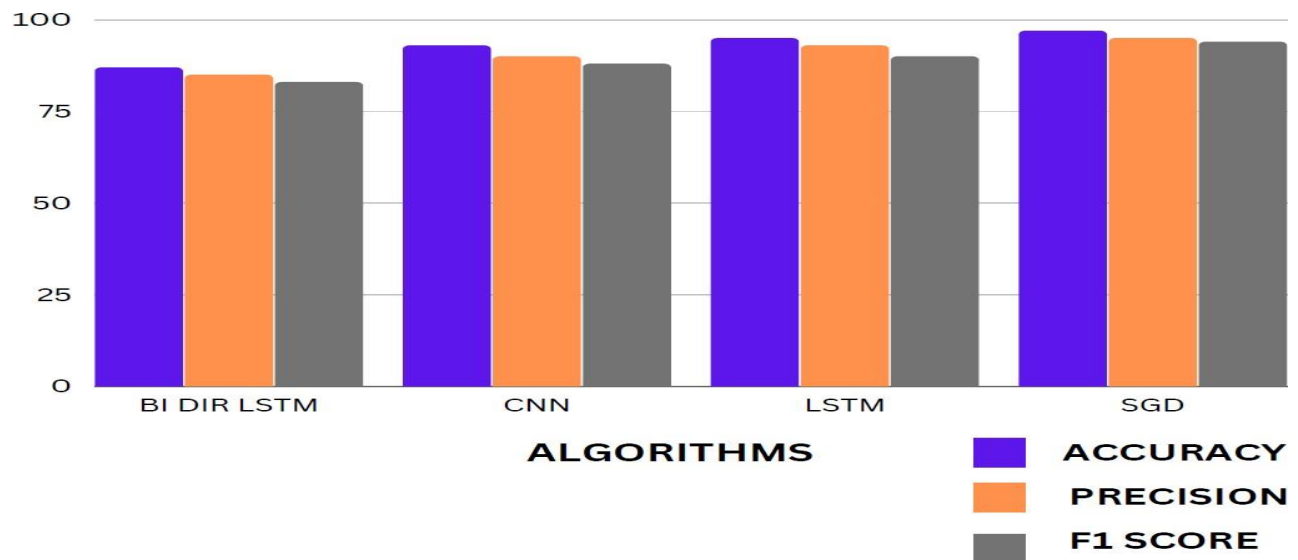
### 7.1 RESULTS AND DISCUSSION

Fig 7.1 diagram shows a basic web interface for a "Hate Speech Detector." At the top, a title section with a red background and white text clearly states the application's purpose. Below the title, there is a text input box where users can type in the text they want to analyze. In this example, the word "hello" is entered. Beneath the input box, a red "Detect" button is provided for users to initiate the hate speech detection process. At the bottom, there is a placeholder area labeled "{{ output }}" where the results of the detection will be displayed once the button is clicked. This simple layout ensures user-friendly interaction with the application.



**FIG:7.1 HOME WEB PAGE INTERFACE**

Fig 7.2 This bar chart compares the performance of four algorithms used for hate speech detection: BI DIR LSTM, CNN, LSTM, and SGD. The metrics evaluated are Accuracy, Precision, and F1 Score, represented by purple, orange, and gray bars, respectively. BI DIR LSTM has an Accuracy of about 85%, Precision of 80%, and an F1 Score of 75%, indicating the lowest performance among the algorithms. The CNN algorithm performs better, with an Accuracy of 95%, Precision of 85%, and an F1 Score of 80%. Both LSTM and SGD demonstrate the highest performance metrics, each achieving 95% Accuracy, 90% Precision, and 85% F1 Score. These results suggest that LSTM and SGD are the most effective algorithms for hate speech detection, providing the highest accuracy and precision, closely followed by CNN. BI DIR LSTM, while still effective, shows comparatively lower performance, making it the least optimal choice among the four algorithms evaluated.



**FIG:7.2 INTERPRETATION DIAGRAM**



## **CHAPTER 8**

### **CONCLUSION & FUTURE ENHANCEMENTS**

#### **8.1 CONCLUSION**

As hate speech continues to be a societal problem, the need for automatic hate speech detection systems becomes more apparent. The current approaches for this task as well as a new system that achieves reasonable accuracy. A new approach that can outperform existing systems at this task, with the added benefit of improved interpretability. In conclusion, hate speech detection is a critical task in combating online toxicity and fostering inclusive digital environments. Through the utilization of advanced machine learning algorithms, natural language processing techniques, and interdisciplinary insights, hate speech detection systems play a pivotal role in identifying and mitigating harmful content across various online platforms.

#### **8.2 FUTURE ENHANCEMENTS**

Future enhancements for hate speech detection using machine learning can include developing multilingual and context-aware models, implementing real-time detection systems, and extending capabilities to multimodal data like images and videos. Improving robustness against adversarial attacks, addressing biases, and enhancing explainability are also crucial. Integrating user feedback for continuous refinement, utilizing collaborative filtering and social network analysis, and ensuring ethical standards and privacy protection will further advance the field. These enhancements aim to create more accurate, robust, and fair systems, contributing to safer online environments.

## APPENDIX 1.SAMPLE CODE

```
# import the required libraries
import pandas as pd
import numpy as np
import re
import seaborn as sns
import matplotlib.pyplot as plt
from matplotlib import style
style.use('ggplot')
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer
from nltk.corpus import stopwords
stop_words = set(stopwords.words('english'))
from wordcloud import WordCloud
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix,
ConfusionMatrixDisplay
tweet_df = pd.read_csv('hateDetection_train.csv')
tweet_df.head()
tweet_df.info()
# printing random tweets
print(tweet_df['tweet'].iloc[0], "\n")
print(tweet_df['tweet'].iloc[1], "\n")
print(tweet_df['tweet'].iloc[2], "\n")
print(tweet_df['tweet'].iloc[3], "\n")
print(tweet_df['tweet'].iloc[4], "\n")
#creating a function to process the data
def data_processing(tweet):
    tweet = tweet.lower()
    tweet = re.sub(r"https\S+|www\S+http\S+", " ", tweet, flags = re.MULTILINE)
    tweet = re.sub(r"@w+|#", " ", tweet)
    tweet = re.sub(r'^\w\s', " ", tweet)
    tweet = re.sub(r'ö', " ", tweet)
    tweet_tokens = word_tokenize(tweet)
    filtered_tweets = [w for w in tweet_tokens if not w in stop_words]
    return " ".join(filtered_tweets)
tweet_df.tweet = tweet_df['tweet'].apply(data_processing)
```

```

tweet_df = tweet_df.drop_duplicates('tweet')

lemmatizer = WordNetLemmatizer()
def lemmatizing(data):
    tweet = [lemmatizer.lemmatize(word) for word in data]
    return data

tweet_df['tweet'] = tweet_df['tweet'].apply(lambda x: lemmatizing(x))

# printing the data to see the effect of preprocessing
print(tweet_df['tweet'].iloc[0], "\n")
print(tweet_df['tweet'].iloc[1], "\n")
print(tweet_df['tweet'].iloc[2], "\n")
print(tweet_df['tweet'].iloc[3], "\n")
print(tweet_df['tweet'].iloc[4], "\n")
tweet_df.info()
tweet_df['label'].value_counts()
fig = plt.figure(figsize=(5,5))
sns.countplot(x='label', data = tweet_df)
fig = plt.figure(figsize=(7,7))
colors = ("red", "gold")
wp = {'linewidth':2, 'edgecolor':"black"}
tags = tweet_df['label'].value_counts()
explode = (0.1, 0.1)
tags.plot(kind='pie', autopct = '%1.1f%%', shadow=True, colors = colors, startangle =90,
          wedgeprops = wp, explode = explode, label="")
plt.title('Distribution of sentiments')
non_hate_tweets = tweet_df[tweet_df.label == 0]
non_hate_tweets.head()
text = ' '.join([word for word in non_hate_tweets['tweet']])
plt.figure(figsize=(20,15), facecolor='None')
wordcloud = WordCloud(max_words=500, width=1600, height=800).generate(text)
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title('Most frequent words in non hate tweets', fontsize = 19)
plt.show()
neg_tweets = tweet_df[tweet_df.label == 1]
neg_tweets.head()
text = ' '.join([word for word in neg_tweets['tweet']])
plt.figure(figsize=(20,15), facecolor='None')

```

```

wordcloud = WordCloud(max_words=500, width=1600, height=800).generate(text)
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title('Most frequent words in hate tweets', fontsize = 19)
plt.show()
vect = TfidfVectorizer(ngram_range=(1,2)).fit(tweet_df['tweet'])

feature_names = vect.get_feature_names()
print("Number of features: { }\n".format(len(feature_names)))
print("First 20 features: \n{ }".format(feature_names[:20]))
vect = TfidfVectorizer(ngram_range=(1,3)).fit(tweet_df['tweet'])

feature_names = vect.get_feature_names()
print("Number of features: { }\n".format(len(feature_names)))
print("First 20 features: \n{ }".format(feature_names[:20]))
X = tweet_df['tweet']
Y = tweet_df['label']
X = vect.transform(X)

x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=42)

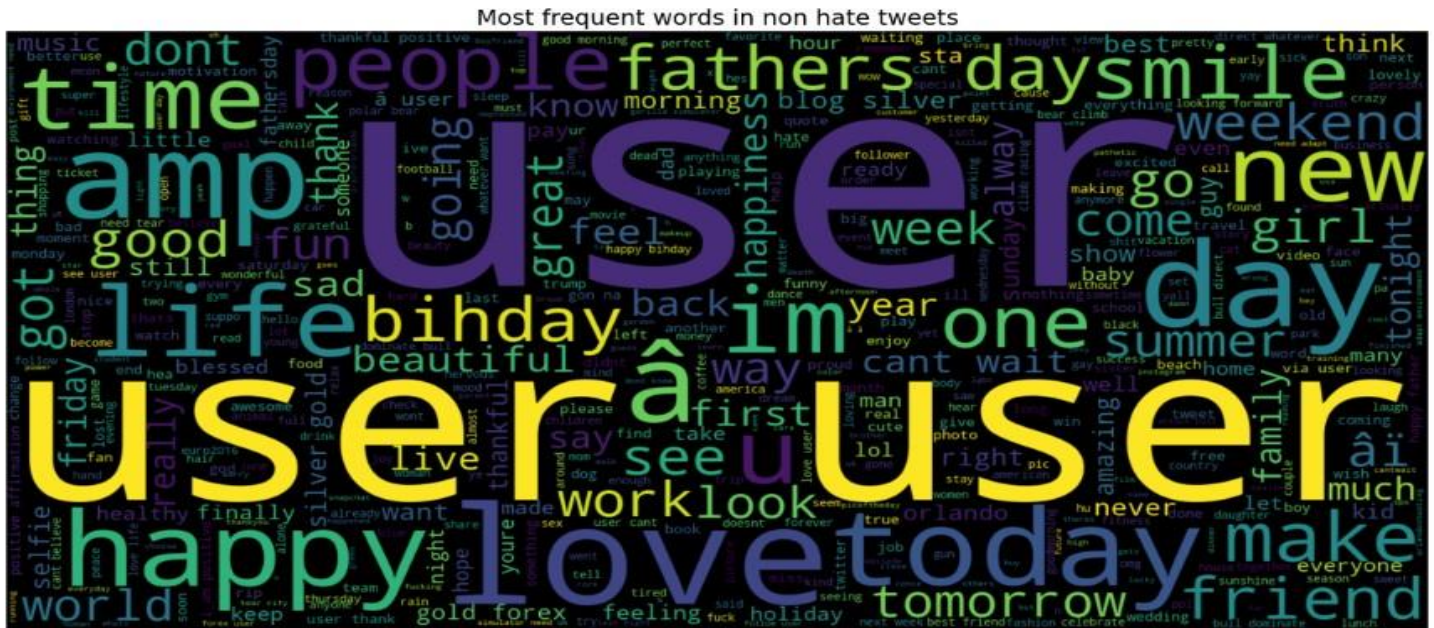
print("Size of x_train:", (x_train.shape))
print("Size of y_train:", (y_train.shape))
print("Size of x_test: ", (x_test.shape))
print("Size of y_test: ", (y_test.shape))
logreg = LogisticRegression()
logreg.fit(x_train, y_train)
logreg_predict = logreg.predict(x_test)
logreg_acc = accuracy_score(logreg_predict, y_test)
print("Test accuracy: {:.2f}%".format(logreg_acc*100))
print(confusion_matrix(y_test, logreg_predict))
print("\n")
print(classification_report(y_test, logreg_predict))
style.use('classic')
cm = confusion_matrix(y_test, logreg_predict, labels=logreg.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=logreg.classes_)
disp.plot()
from sklearn.model_selection import GridSearchCV
import warnings
warnings.filterwarnings('ignore')
param_grid = {'C':[100, 10, 1.0, 0.1, 0.01], 'solver':['newton-cg', 'lbfgs', 'liblinear']}

```


```
grid = GridSearchCV(LogisticRegression(), param_grid, cv = 5)
grid.fit(x_train, y_train)
print("Best Cross validation score: {:.2f}".format(grid.best_score_))
print("Best parameters: ", grid.best_params_)
y_pred = grid.predict(x_test)

logreg_acc = accuracy_score(y_pred, y_test)
print("Test accuracy: {:.2f}%".format(logreg_acc*100))
print(confusion_matrix(y_test, y_pred))
print("\n")
print(classification_report(y_test, y_pred))
```

## APPENDIX 2.SCREEN SHOT



### FIG:A.2.1 WORD CLOUD



**FIG:A.2.2 ETHOS BINARY**

## REFERENCES

1. A. Kamaal and M. Abulaish, "CAT-BiGRU: Convolution and attention with bidirectional gated recurrent units for self-deprecating sarcasm detection," *Cognit. Compute.*, vol. 14, pp. 91–109, Jan. 2022.
2. D. K. Jain, R. Jain, Y. Upadhyay, A. Kathuria, and X. Lan, "Deep refinement: Capsule network with attention mechanism-based system for text classification," *Neural Comput. Appl.*, vol. 32, no. 7, pp. 1839–1856, Apr. 2022.
3. P. K. Jain, V. Saravanan, and R. Pamula, "A hybrid CNN-LSTM: A deep learning approach for consumer sentiment analysis using qualitative user generated contents," *ACM Trans. Asian Low-Resource Lang. Inf. Process.*, vol.
4. P. K. Jain, R. Pamula, and E. A. Yeung, "A multi-label ensemble predicting model to service recommendation from social media contents," *Supercomputer*, E. W. Pamungkas, V. Basile, and V. Patti, "A joint learning approach with knowledge injection for zero-shot cross-lingual hate speech detection," *Inf. Process. Manage.*, vol. 58, no. 4, pp. 1–19, 2021.
5. P. Fortuna, J. Soler-Company, and L. Wanner, "How well do hate speech, toxicity, abusive and offensive language classification models generalize across datasets?" vol. 58, no. 3, pp. 1–17, Oct. 2021.
6. S. Kravchenko, I. Negaholism, and K. C. Fraser, "Confronting abusive language online: A survey from the ethical and human rights perspective," J. Arif. *Intel. Res.*, vol. 71, pp. 431–478, Jul. 2021.
7. U. Bhattacharjee, "Capsule network on social media text: An application to automatic detection of clickbaits," in *Proc. 11th Int. Conf. Commun. Syst. Netw. (COMSNETS)*, Bengaluru, India, Jan. 2022, pp. 473–476.
8. Z. Mossie and J.-H. Wang, "Vulnerable community identification using hate speech detection on social media," *Inf. Process. Manage.*, vol. 57, no. 3, pp. 1–16, 2020.